

# Walmart Stores Sales Forecasting

One challenge of modeling retail data is the need to make decisions based on limited history. If Christmas comes but once a year, so does the chance to see how strategic decisions impacted the bottom line.

In this use case, provided with historical sales data for 45 Walmart stores located in different regions. Each store contains many departments, and participants must project the sales for each department in each store.

In addition, Walmart runs several promotional markdown events throughout the year. These markdowns precede prominent holidays, the four largest of which are the Super Bowl, Labor Day, Thanksgiving, and Christmas. The weeks including these holidays are weighted five times higher in the evaluation than non-holiday weeks. Part of the challenge presented by this competition is modeling the effects of markdowns on these holiday weeks in the absence of complete/ideal historical data.

## ML Problem Formulation

### Time-series forecasting and Regression

- To forecast sales for 45 stores, given sales data for 45 Walmart stores located in different regions and time.

To solve the above we would be using data collected in Jan - Mar 2015 to predict the pickups in Jan - Mar 2016.

2010-02-05 to 2012-11-01

### Performance Metric

This competition is evaluated on the weighted mean absolute error (WMAE):

$WMAE = \frac{1}{\sum w_i} \sum_{i=1}^n w_i |y_i - \hat{y}_i|$  where

$n$  is the number of rows  $\hat{y}_i$  is the predicted sales  $y_i$  is the actual sales  $w_i$  are weights.  $w = 5$  if the week is a holiday week, 1 otherwise

### Data

#### stores.csv

This file contains anonymized information about the 45 stores, indicating the type and size of store.

**train.csv**

This is the historical training data, which covers to 2010-02-05 to 2012-11-01. Within this file you will find the following fields:

Store - the store number Dept - the department number Date - the week Weekly\_Sales - sales for the given department in the given store IsHoliday - whether the week is a special holiday week  
test.csv

This file is identical to train.csv, except we have withheld the weekly sales. You must predict the sales for each triplet of store, department, and date in this file.

**features.csv**

This file contains additional data related to the store, department, and regional activity for the given dates. It contains the following fields:

Store - the store number Date - the week Temperature - average temperature in the region Fuel\_Price - cost of fuel in the region Markdown1-5 - anonymized data related to promotional markdowns that Walmart is running. Markdown data is only available after Nov 2011, and is not available for all stores all the time. Any missing value is marked with an NA. CPI - the consumer price index Unemployment - the unemployment rate IsHoliday - whether the week is a special holiday week For convenience, the four holidays fall within the following weeks in the dataset (not all holidays are in the data):

Super Bowl: 12-Feb-10, 11-Feb-11, 10-Feb-12, 8-Feb-13 Labor Day: 10-Sep-10, 9-Sep-11, 7-Sep-12, 6-Sep-13 Thanksgiving: 26-Nov-10, 25-Nov-11, 23-Nov-12, 29-Nov-13 Christmas: 31-Dec-10, 30-Dec-11, 28-Dec-12, 27-Dec-13

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
```

```
In [2]: train_data = pd.read_csv('train.csv')
train_data.head()
```

Out[2]:

	Store	Dept	Date	Weekly_Sales	IsHoliday
0	1	1	2010-02-05	24924.50	False
1	1	1	2010-02-12	46039.49	True
2	1	1	2010-02-19	41595.55	False
3	1	1	2010-02-26	19403.54	False
4	1	1	2010-03-05	21827.90	False

In [3]: `train_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 421570 entries, 0 to 421569
Data columns (total 5 columns):
Store          421570 non-null int64
Dept           421570 non-null int64
Date           421570 non-null object
Weekly_Sales   421570 non-null float64
IsHoliday      421570 non-null bool
dtypes: bool(1), float64(1), int64(2), object(1)
memory usage: 13.3+ MB
```

In [4]: `store_data = pd.read_csv('stores.csv')`  
`store_data.head()`

Out[4]:

	Store	Type	Size
0	1	A	151315
1	2	A	202307
2	3	B	37392
3	4	A	205863
4	5	B	34875

In [5]: `features_data = pd.read_csv('features.csv')`  
`features_data.head()`

Out[5]:

	Store	Date	Temperature	Fuel_Price	MarkDown1	MarkDown2	MarkDown3	MarkDown4	MarkI
0	1	2010-02-05	42.31	2.572	NaN	NaN	NaN	NaN	
1	1	2010-02-12	38.51	2.548	NaN	NaN	NaN	NaN	
2	1	2010-02-19	39.93	2.514	NaN	NaN	NaN	NaN	
3	1	2010-02-26	46.63	2.561	NaN	NaN	NaN	NaN	
4	1	2010-03-05	46.50	2.625	NaN	NaN	NaN	NaN	

```
In [6]: dataset = train_data.merge(features_data, how='left').merge(store_data, how='left')
dataset.head()
```

Out[6]:

	Store	Dept	Date	Weekly_Sales	IsHoliday	Temperature	Fuel_Price	MarkDown1	MarkDown2
0	1	1	2010-02-05	24924.50	False	42.31	2.572	NaN	NaN
1	1	1	2010-02-12	46039.49	True	38.51	2.548	NaN	NaN
2	1	1	2010-02-19	41595.55	False	39.93	2.514	NaN	NaN
3	1	1	2010-02-26	19403.54	False	46.63	2.561	NaN	NaN
4	1	1	2010-03-05	21827.90	False	46.50	2.625	NaN	NaN

```
In [7]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 421570 entries, 0 to 421569
Data columns (total 16 columns):
Store          421570 non-null int64
Dept           421570 non-null int64
Date           421570 non-null object
Weekly_Sales   421570 non-null float64
IsHoliday      421570 non-null bool
Temperature    421570 non-null float64
Fuel_Price     421570 non-null float64
MarkDown1      150681 non-null float64
MarkDown2      111248 non-null float64
MarkDown3      137091 non-null float64
MarkDown4      134967 non-null float64
MarkDown5      151432 non-null float64
CPI            421570 non-null float64
Unemployment   421570 non-null float64
Type           421570 non-null object
Size           421570 non-null int64
dtypes: bool(1), float64(10), int64(3), object(2)
memory usage: 51.9+ MB
```

## Data Cleaning/ Preprocessing

```
In [8]: # Handling missing values
```

```
dataset['MarkDown1'].fillna(0, inplace=True)
dataset['MarkDown2'].fillna(0, inplace=True)
dataset['MarkDown3'].fillna(0, inplace=True)
dataset['MarkDown4'].fillna(0, inplace=True)
dataset['MarkDown5'].fillna(0, inplace=True)
```


## Exploratory Data Analysis

```
In [9]: # There are 1285 records with negative Weekly_Sales which means refund.  
  
dataset['Weekly_Sales'][dataset['Weekly_Sales']<0].count()  
  
# Kept negative values as is, as its giving bad result after replacing those
```

Out[9]: 1285


```
In [10]: # We have very few records with Holidays  
  
sns.countplot(dataset['IsHoliday'])
```

Out[10]: <matplotlib.axes.\_subplots.AxesSubplot at 0x5683924da0>






```
In [11]: # In data most of the store are of Type 'A' and very few of type 'C'  
  
sns.countplot(dataset['Type'])
```




Out[11]: <matplotlib.axes.\_subplots.AxesSubplot at 0x56838ab2b0>



```
In [12]: # Function to plot scatter plot between selected feature and Weekly_Sales
def scatter(dataset, column):
    plt.figure()
    plt.scatter(dataset[column], dataset['Weekly_Sales'])
    plt.ylabel('Weekly_Sales')
    plt.xlabel(column)
```

```
In [13]: scatter(dataset, 'Fuel_Price')
scatter(dataset, 'Size')
scatter(dataset, 'CPI')
scatter(dataset, 'Temperature')
scatter(dataset, 'Store')
scatter(dataset, 'Dept')
```



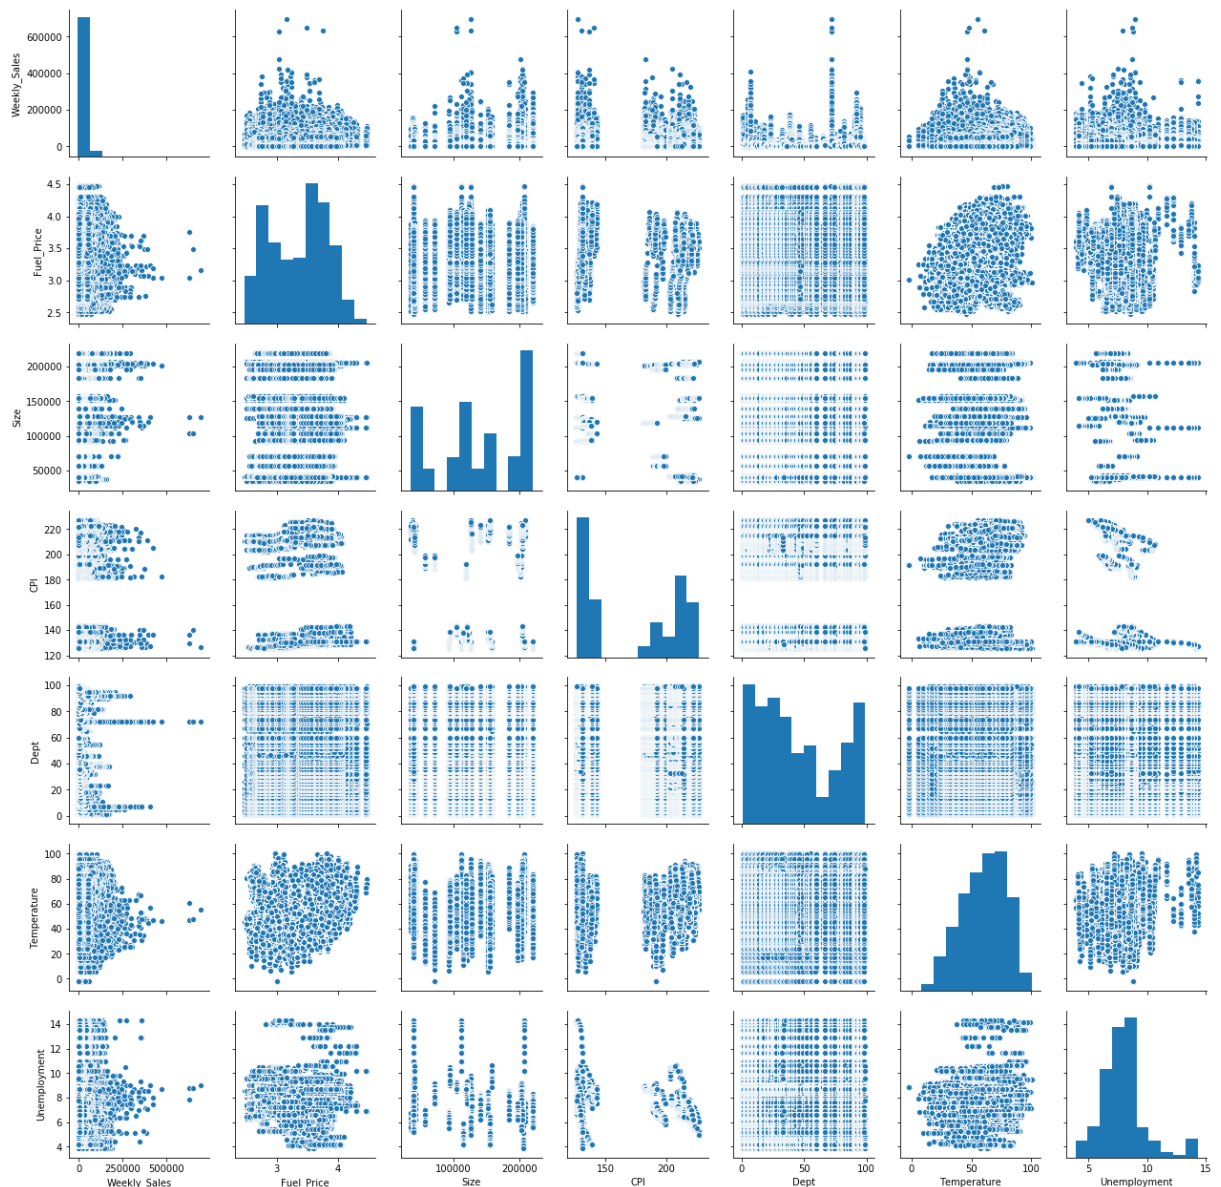


- Fuel\_Price, Temperature and Size feature seems to be interesting








```
In [14]: sns.pairplot(dataset, vars=['Weekly_Sales', 'Fuel_Price', 'Size', 'CPI', 'Dept', 'Temperature', 'Unemployment'])
```

```
Out[14]: <seaborn.axisgrid.PairGrid at 0x5686389550>
```



```
In [15]: scatter(dataset, 'MarkDown1')  
scatter(dataset, 'MarkDown2')  
scatter(dataset, 'MarkDown3')  
scatter(dataset, 'MarkDown4')  
scatter(dataset, 'MarkDown5')
```





- Markdown3 seems interesting, others Markdown doesn't seem to be an important feature

## Creating new features

```
In [16]: # Formatting the date
dataset['Date'] = pd.to_datetime(dataset['Date'], format='%Y-%m-%d')

# Creating new feature from date i.e. Year and week
dataset['Year'] = dataset['Date'].dt.year
dataset['Week'] = dataset['Date'].dt.week

# dataset['Month'] = dataset['Date'].dt.month

# Map values of IsHoliday with 0 and 5 as per weights given to them
dataset['IsHoliday'] = dataset['IsHoliday'].map({False:0, True:5})
```

In [17]: *# Dropping not useful features*

```
dataset = dataset.drop(['Date', 'MarkDown1', 'MarkDown2', 'MarkDown4', 'MarkDown5', 'Fuel_Price'], axis=1)
dataset.head()
```


Out[17]:

	Store	Dept	Weekly_Sales	IsHoliday	Temperature	MarkDown3	Size	Year	Week
0	1	1	24924.50	0	42.31	0.0	151315	2010	5
1	1	1	46039.49	5	38.51	0.0	151315	2010	6
2	1	1	41595.55	0	39.93	0.0	151315	2010	7
3	1	1	19403.54	0	46.63	0.0	151315	2010	8
4	1	1	21827.90	0	46.50	0.0	151315	2010	9

In [20]: *# Plotting correlation between all important features*

```
corr = dataset.corr()
plt.figure(figsize=(15, 10))
sns.heatmap(corr, annot=True)
plt.plot()
```

Out[20]: []



**Split of data for train and test**

```
In [21]: # Used 2010 and 2011 data for training
# Used 2012 data for test

train_x = dataset.loc[dataset['Year']<2012]
train_y = dataset['Weekly_Sales'].loc[dataset['Year']<2012]
train_x = train_x.drop('Weekly_Sales', axis=1)

print("Records in training data :", train_x.shape)
print("Records in training target label :", train_y.shape)

test_x = dataset.loc[dataset['Year']>2011]
test_y = dataset['Weekly_Sales'].loc[dataset['Year']>2011]
test_x = test_x.drop('Weekly_Sales', axis=1)

print("Records in test data :", test_x.shape)
print("Records in test target label :", test_y.shape)
```

```
Records in training data : (294132, 8)
Records in training target label : (294132,)
Records in test data : (127438, 8)
Records in test target label : (127438,)
```

## Modeling

```
In [22]: def calculate_error(test_y, predicted, weights):
return mean_absolute_error(test_y, predicted, sample_weight=weights)
```

```
In [23]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error
from sklearn.model_selection import RandomizedSearchCV
```

## Linear Regression

```
In [24]: def plots(test_y, pred_y):
plt.figure(figsize=(15,15))

plt.subplot(231)
sns.set_style('whitegrid')
sns.kdeplot(np.array(test_y), bw=0.5)
plt.xlabel('Actual Value')

plt.subplot(232)
sns.kdeplot(np.array(pred_y), bw=0.5)
plt.xlabel('Predicted Value')

plt.subplot(233)
delta_y = test_y - pred_y;
sns.kdeplot(np.array(delta_y), bw=0.5)
plt.xlabel('Errors')
plt.show()
```


```
In [25]: model = LinearRegression()
model.fit(train_x, train_y)

pred_y = model.predict(test_x)

weights = test_x['IsHoliday']
LR_error = calculate_error(test_y, pred_y, weights)
print("Error: ", LR_error)
```

Error: 14744.732702185725

```
In [26]: # Plot between Actual Value vs Predicted Value
plots(test_y, pred_y)
```



- Actual values and predicted values are very different. Hence error is very high

## RandomForestRegressor

```
In [27]: from sklearn.ensemble import RandomForestRegressor
```

### *Hyperparamter tuning*

```
In [28]: n_estimators = [50, 100, 150, 200]
max_depth=[5, 10, 15, 20]
min_samples_split = [2, 4, 6, 8]
params = {'n_estimators': n_estimators, 'max_depth': max_depth, 'min_samples_split': min_samples_split}

randomForest = RandomForestRegressor()
randomCV = RandomizedSearchCV(randomForest, params, cv=5)
randomCV.fit(train_x, train_y)

optimal_n_estimators = randomCV.best_params_['n_estimators']
optimal_max_depth = randomCV.best_params_['max_depth']
optimal_min_samples_split = randomCV.best_params_['min_samples_split']

print("Best parameters: ", randomCV.best_params_)
print("Best estimators: ", randomCV.best_estimator_)
```

```
Best parameters: {'n_estimators': 200, 'min_samples_split': 8, 'max_depth': 10}
Best estimators: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=10,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=8,
min_weight_fraction_leaf=0.0, n_estimators=200, n_jobs=None,
oob_score=False, random_state=None, verbose=0, warm_start=False)
```


```
In [29]: rfModel = RandomForestRegressor(n_estimators=optimal_n_estimators, max_depth=optimal_max_depth,
min_samples_split=optimal_min_samples_split)
rfModel.fit(train_x, train_y)

pred_y = rfModel.predict(test_x)

weights = test_x['IsHoliday']
RF_error = calculate_error(test_y, pred_y, weights)
print("Error: ", RF_error)
```

```
Error: 4377.498592401463
```

```
In [30]: plots(test_y, pred_y)
```



- Actual values and predicted values are better than Linear Regression model and error is also less

## ExtraTreesRegressor

### Hyper parameter tuning

```
In [31]: from sklearn.ensemble import ExtraTreesRegressor
```



```
In [32]: n_estimators = [50, 100, 150, 200]
max_depth=[5, 10, 15, 20]
min_samples_split = [2,4 ,8, 10]
params = {'n_estimators' : n_estimators, 'max_depth': max_depth, 'min_samples_spli

etree = ExtraTreesRegressor()
randomCV = RandomizedSearchCV(etree, params , cv=5)
randomCV.fit(train_x, train_y)

optimal_n_estimators = randomCV.best_params_['n_estimators']
optimal_max_depth = randomCV.best_params_['max_depth']
optimal_min_samples_split = randomCV.best_params_['min_samples_split']

print("Best parameters: ", randomCV.best_params_)
print("Best estimators: ", randomCV.best_estimator_)
```

Best parameters: {'n\_estimators': 150, 'min\_samples\_split': 8, 'max\_depth': 20}

Best estimators: ExtraTreesRegressor(bootstrap=False, criterion='mse', max\_depth=20,

max\_features='auto', max\_leaf\_nodes=None,  
min\_impurity\_decrease=0.0, min\_impurity\_split=None,  
min\_samples\_leaf=1, min\_samples\_split=8,  
min\_weight\_fraction\_leaf=0.0, n\_estimators=150, n\_jobs=None,  
oob\_score=False, random\_state=None, verbose=0, warm\_start=False)


```
In [33]: extraTree = ExtraTreesRegressor(n_estimators=optimal_n_estimators, max_depth = opt
min_samples_split = optimal_min_samples_split, n_j
extraTree.fit(train_x, train_y)

pred_y = extraTree.predict(test_x)

weights = test_x['IsHoliday']
ET_error = calculate_error(test_y, pred_y, weights)
print("Error: ", ET_error)
```

Error: 1981.3401446754774

```
In [34]: plots(test_y, pred_y)
```



- Actual value and predicted value seems very close now and results are much better using ExtraTrees Model

## XGBRegressor

```
In [35]: from xgboost import XGBRegressor
```

```
In [36]: xgRegressor = XGBRegressor(learning_rate=0.1, n_estimators=100, max_depth=3, min_
        gamma=0, subsample=0.8, reg_alpha=200, reg_lambda=200, colsample_byt


xgRegressor.fit(train_x, train_y)

pred_y = xgRegressor.predict(test_x)

weights = test_x['IsHoliday']
XG_error = calculate_error(test_y, pred_y, weights)
print("Error: ", XG_error)
```

Error: 7310.941025042182

In [37]: `plots(test_y, pred_y)`



## Stacking

In [38]: `from mlxtend.regressor import StackingRegressor`

```
In [39]: lrReg = LinearRegression()

rfReg = RandomForestRegressor(n_estimators=100, max_depth=10,
                             min_samples_split=2)
xgbReg = XGBRegressor(learning_rate=0.1, n_estimators=100, max_depth=3, min_child
                      gamma=0, subsample=0.8, reg_alpha=200, reg_lambda=200, colsample_byt
xtraTreesReg = ExtraTreesRegressor(n_estimators=150, max_depth = 15, min_samples_s


stackRgr = StackingRegressor(regressors=[lrReg, rfReg, xgbReg], meta_regressor=xtr

# Training the stacking classifier
stackRgr.fit(train_x, train_y)
stack_pred_y = stackRgr.predict(test_x)

weights = test_x['IsHoliday']
Stack_error = calculate_error(test_y, stack_pred_y, weights)
print("Error: ", Stack_error)
```

Error: 3189.9629174081642

```
In [40]: plots(test_y, stack_pred_y)
```



## Test Data for submission

```
In [41]: test_data = pd.read_csv('test.csv')
test_data.head()
```

Out[41]:

	Store	Dept	Date	IsHoliday
0	1	1	2012-11-02	False
1	1	1	2012-11-09	False
2	1	1	2012-11-16	False
3	1	1	2012-11-23	True
4	1	1	2012-11-30	False

```
In [42]: # Merge test data with features data
test_data = test_data.merge(features_data, how='left').merge(store_data, how='left')
```

```
In [43]: # handle missing values
test_data['Markdown3'] = test_data['Markdown3'].fillna(0)
```

```
In [44]: # Formating the date
test_data['Date'] = pd.to_datetime(test_data['Date'], format='%Y-%m-%d')

# Creating new features year and week from date
test_data['Year'] = test_data['Date'].dt.year
test_data['Week'] = test_data['Date'].dt.week

# test_data['Month'] = test_data['Date'].dt.month

test_data['IsHoliday'] = test_data['IsHoliday'].map({False:0, True:5})
```

```
In [45]: # Dropping not important features
final_testdata = test_data.drop(['Date', 'MarkDown1', 'MarkDown2', 'MarkDown4', 'Mark
                                'CPI', 'Fuel_Price'], axis=1)
final_testdata.head()
```

Out[45]:

	Store	Dept	IsHoliday	Temperature	MarkDown3	Size	Year	Week
0	1	1	0	55.32	50.82	151315	2012	44
1	1	1	0	61.24	40.28	151315	2012	45
2	1	1	0	52.92	103.78	151315	2012	46
3	1	1	5	56.23	74910.32	151315	2012	47
4	1	1	0	52.34	3838.35	151315	2012	48

```
In [46]: # Using ExtraTress Model to predict result for test.csv
pred_test_y = extraTree.predict(final_testdata)

final_testdata['Weekly_Sales'] = pred_test_y
final_testdata['Id'] = test_data['Store'].astype(str) + '_' + test_data['Dept'].ast

test_result = final_testdata[['Id', 'Weekly_Sales']]
test_result.to_csv('Test_results.csv', index=False)
```

```
In [47]: from prettytable import PrettyTable

table = PrettyTable()
table.add_column('S.No.', [1, 2, 3, 4, 5])
table.add_column('Model', ['Linear Regression', 'Random Forest Regression', 'Extra T
                        'XgBoost Regression', 'Stacking'])
table.add_column('Train Error', [LR_error, RF_error, ET_error, XG_error, Stack_err
print(table)
```

S.No.	Model	Train Error
1	Linear Regression	14744.732702185725
2	Random Forest Regression	4377.498592401463
3	Extra Trees Regression	1981.3401446754774
4	XgBoost Regression	7310.941025042182
5	Stacking	3189.9629174081642

## Conclusion

Steps Followed:

- Data Cleaned - Replaced NaN values
- Performed EDA and found out important features
- Created new features from Date like year and week
- Created various models using LinearRegression, RandomForestRegressor, ExtraTressRegressor and XGBoostRegressor

- Also created model using stacking
- Plotted Actual values VS predicted values

**Among 5 models, ExtraTressRegressor gives the best result**